


Generated Code

What a large language model is; a generated program as notation;
the compiler, the test and the model checker as the gate.

← CC 12 · Bounded Model Checking

WHAT IT IS

A large language model is a machine trained on **notation**, asked for **meaning**.



tokens in · a distribution over the next one out · then a sample


Context in, a probability for every next token out, one sampled and appended. Trained on text by predicting the next token: a purely syntactic signal. Meaning never handed to it.

It is arithmetic on bits — in principle a C* program, which mipster would run. It computes what a Turing machine computes, and week 8 said what that means about self-certification.

Its behaviour on every context is one of week 7's infinite sheets. Only sampled, never enumerated. Testing shows presence.

HALLUCINATION

A hallucination is a **proof-shaped object** that isn't **true**.



plausible · provable · true

Fluency is syntax. Correctness is semantics. A generated program can be perfectly fluent C — indented, idiomatic, commented — and not parse as C*, or parse and divide by zero, or run and compute the wrong thing. Three tests, and they are not the same test.

Not a defect to patch away: the gap of week 7 at consumer scale. The response is the same as everywhere in this class: check the notation against a semantics you own.

Ask a machine for a C* program. Then *starc*, *mipster*, *rotor*.


```
$ ./selfie -c generated.c
./selfie: syntax error in generated.c in line 3: unexpected
// ... after asking again for C*, not C:
$ ./selfie -c generated.c -m 1 7
5040
$ ./rotor -c generated.c - 0 && tools/bitme.py -kmax 100
... core-0-division-by-zero - satisfiable · input 0
```

Does it parse? The compiler decides, always.
Usually not on the first try: the machine produces C fluently, and C* is not C.

Does it do what was asked? Run it on the inputs you thought of. Presence, not absence.

Can it fail? Within a thousand steps, on any input: rotor and bitme. Absence, within the bound. None of the three is available to the machine that wrote the program, because the only semantics in the loop is the one you brought.

A generator supplies notation. The semantics has to come from **somewhere else**.



the only part the machine does not supply

what the machine supplies · and what it does not

Cantor's list needed an object built from outside it. Gödel's system needed a stronger one. Thompson's compiler needed a second compiler. The generator needs a check it cannot be. One theorem, four costumes, and this class met all four.

Models grading models, models trained on model output, agents invoking themselves: the loops are closed, and week 8 said what no loop establishes about itself. Spend the effort on the outside check.

WHAT CHANGED

Generation got **cheap**. Verification did **not**. This class was about verification.

Week 11: hard to find, easy to check. That is a description of a healthy relationship with a machine: accept work you could not have produced, provided you can check it. When you cannot, delegation is faith.

Value migrates to the two ends the machine does not occupy: stating what should be true, and establishing that it is. A grammar, a type, a test, a bad state and a bound. Every one of those is something you wrote this semester.

FOR THE CAPSTONE

Use a machine to draft struct-execution if you like. Then run the three checks, and note in your submission which of them found what. That note is worth more than the draft.

Listen, then read.

LISTEN

[Strauss · Capriccio](#) — Runnicles, San Francisco, Kiri Te Kanawa. His last opera argues for an evening about which comes first, the words or the music: notation or meaning. The Countess never chooses, and the gate is the reason she cannot.

READ

[The Alignment Problem](#) — Christian, 2020: what it means to get a machine to do what was wanted, told through the people who tried. And, more technical, [Evaluating Large Language Models Trained on Code](#) — Chen and others, 2021: the Codex paper, with the honest table of how often the generated program passes the tests.