

Everything is bits.

Numbers, negative numbers, overflow, characters, text, files,
images, code — one notation, and what the machine makes of it.

← ICS 02 · Size

BINARY


85 is **1010101**. Nothing else changed.

Decimal: $85 = 8 \cdot 10 + 5 \cdot 1$. Binary: $1010101 = 1 \cdot 64 + 0 \cdot 32 + 1 \cdot 16 + 0 \cdot 8 + 1 \cdot 4 + 0 \cdot 2 + 1 \cdot 1$.


Same number, a different base. The value is the meaning; the digits are the notation.

Adding works the same way in any base — carry when a column overflows its base. $1010101 + 111 = 1011100$, which is $85 + 7 = 92$.

Hexadecimal groups four bits into one digit: 1010101 is 0x55. Octal groups three. Both are just shorter ways to write the same bits; the machine only ever sees the bits.



Above base one: **exponentially** shorter.
 Above base two: only a **constant factor**.



Unary writes n as n bars; positional notation needs $\lceil \log_b n \rceil$ digits: 85 is 85 bars, 7 bits, 2 decimal digits. Base 1 to 2 buys an exponential; base 2 to 10 buys $\log_2 10 \approx 3.3$, for every n .

The price is arithmetic. In unary, counting is appending a bar: constant time. In any positional base an addition walks the digits: $\log n$ steps, with carries.

So the mathematics says leave unary, and does not say which base. That is a decision for physics.

A transistor is a **switch**. Two levels leave the most room for noise, and nothing to calibrate.

B LEVELS PER DIGIT

A base- b digit must hold one of b distinguishable states. The gap between neighbouring levels is what noise has to cross, and it shrinks with b . Two levels, on and off, put that gap as far apart as the supply voltage allows.

COST PER LEVEL, NOT PER DIGIT

Hardware pays per level, not per digit. If a digit costs b and you need $\log_b n$ of them, $b / \ln b$ is smallest near e : 3 beats 2 by 5 percent. Five percent buys no factory; a switch that needs no calibration does.

WHERE HIGHER BASES DO WIN

Flash memory stores 2 to 4 bits per cell as 4 to 16 charge levels: a constant-factor saving, paid for in error correction, and read back out as bits.

And DNA uses base 4. Its digits are molecules, not voltage levels: four distinct shapes, each pairing with exactly one other, so four states cost nothing in noise margin, and pairing is how a strand is copied and checked. A different physics, a different base.

Boolean algebra lets the rest of the machine forget the physics: two values, and every gate a function of them.

Balanced ternary: digits $-1, 0, 1$. Negation is a **digit flip**, and no sign bit is needed.

Write base 3 with the digits $-1, 0$ and 1 instead of $0, 1$ and 2 . Then every integer, negative ones included, has exactly one representation, and $-n$ is n with every digit negated: $85 = 81 + 3 + 1$ is $1\ 0\ 0\ 1\ 1$ and -85 is $-1\ 0\ 0\ -1\ -1$. No two's complement, no sign convention, and rounding is truncation.


Setun, built at Moscow State University in 1958, computed in balanced ternary and about fifty were made. Knuth called it the prettiest number system of all. It lost anyway, and not to a better idea: to the transistor, which is a switch, and to a supply chain of parts that were all binary.

	BINARY	BALANCED TERNARY
digits	0 1	-1 0 1
negative numbers	a convention on top	built in
negation	invert and add 1	flip every digit
digit count for n	$\log_2 n$	$\log_3 n = 0.63 \cdot \log_2 n$
levels a digit holds	2	3

Mathematically the nicer system; physically three levels where two would do. The bits do not say which; the transistor does.

NEGATIVE NUMBERS

Two digits, read three ways.



With two decimal digits you can encode a hundred things. Read them as 0 to 99, or read 50 to 99 as -50 to -1 . Same digits. Subtraction becomes addition: $7 - 8$ is $7 + 92 = 99$, which is -1 .

The machine does exactly this in base two: two's complement. With 64 bits, 2^{64} encodings, read as 0 to $2^{64}-1$ or as -2^{63} to $2^{63}-1$. The bits do not say which.

1010101 with seven bits is 85 unsigned — and -43 signed. Same seven bits.

Finite means the numbers wrap around.

With 64 bits, $2^{64} - 1$ plus 1 is 0. Nothing crashes; the carry out of the top bit is simply lost. That is an overflow, and the machine does not tell you.

Signed, the same wrap makes $2^{63} - 1$ plus 1 equal to -2^{63} : the largest positive number plus one is the most negative one.

Every integer in every program you will ever run lives in a finite box. Most of the time nobody notices. The times somebody does are famous.

```
$ cat examples/overflows.c
// ... prints UINT64_MAX + 1, INT64_MAX + 1, and
$ ./selfie -c examples/overflows.c -m 1
...
```

1010101 is also the letter **U**. People in the 1960s sat down and agreed.

ASCII: seven bits, 128 characters, a table agreed in 1963. 85 is U. 48 is the digit 0, 65 is A, 97 is a, 32 is a space. The digit '0' and the number 0 are different things with different codes.


Unicode extends the table to every script, and UTF-8 encodes it in one to four bytes so that ASCII is unchanged. A byte is eight bits, a nibble four, and the figure names its ends.



Memory is bytes with **addresses**. An address is a number. So a byte can hold an address.




storage and addresses



a value read as an address: a pointer

A string is bytes in a row, ended by a **zero**. A file is a string with a name.




Contiguous: the letters of a word one after the other. The zero byte says where it ends, which is how a machine that only sees bytes knows a word is over.


Non-contiguous: a text is paragraphs, each contiguous, reached through pointers. A directory is names and pointers to files. A file system is a tree of directories, and a pathname is the route from the root.

All of it is bytes at addresses. The structure is in how they are read.

A picture is numbers in rows. A film is pictures in a row. A sound is numbers in a row.




an image, row-major



audio, 8-bit samples

And an **instruction** is 32 bits. Code is bytes too.



Selfie compiles its source to 43,492 instructions of 32 bits each: 173,968 bytes of code and 14,424 bytes of data, laid out one after the other in a file.

That file is bytes at addresses, like a text or an image. The machine reads it and does what it says. Nothing in the bytes tells you whether they are code or data; the machine chapter shows what does.

```
$ ./selfie -c examples/hello-world.c -m 1
./selfie: selfie compiling examples/hello-world.c to 64-
...
./selfie: 64-bit mipster executing 64-bit RISC-U binary
Hello World!
```

1010101 is 85, and U, and -43 , and an address, and half an instruction. The bits do not say which.

Everything a machine stores is a number.
Everything a number means is decided by
whoever reads it, and by nothing else.

Last week: one prefix, two meanings. This
week: one byte, five. The gap between the
notation and its meaning is not a defect of
computers. It is what this class is about, and
next week we start on the notations that pin
meaning down.

*Whatever you see on a screen could
come from anywhere and mean
anything.*

THE BOOK, LIFE 4

Recommended exercises.

- 1 Read the Size chapter from *Numbers to Code*.
- 2 Write 42, 255, and 1000 in binary and hexadecimal. Add 85 and 42 in binary with carries.
- 3 How many digits does 1000 need in unary, binary, ternary and decimal? Which of the four ratios stays the same as the number grows, and which does not?
- 4 What is 1010101 as a signed 7-bit number? What is 11111111 as an unsigned and as a signed byte?
- 5 Write your first name in ASCII, in binary, and count the bits. Then in UTF-8 if it has an umlaut.
- 6 Run `examples/overflows.c` and explain each line of output.
- 7 In the pointers figure, what happens if the byte at address 0 held 7 instead of 85?

NEXT WEEK

Notation: formal languages. EBNF, and the two languages selfie is made of — C* with seven keywords, RISC-U with fourteen instructions — read exactly.

Listen, then read.

LISTEN

[Mozart · Symphony No. 40 in G minor](#) — Harnoncourt, Concentus Musicus. One two-note sigh, the smallest unit, and the whole first movement is that figure at different positions. Everything is bits, in G minor.

READ

[Code: The Hidden Language of Computer Hardware and Software](#) — Petzold, second edition 2022: from flashlights and Morse code to a working computer, one bit at a time. And, more technical, [A Mathematical Theory of Communication](#) — Shannon, 1948: where the bit got its name.